



DMIF University of Udine

Managing Time Series with MongoDB

Paolo Gallo
paolo.gallo@uniud.it

May 18, 2020



- 1 Introduction
- 2 Schema design(s)
- 3 Impact on requirements and performances



The United States is beginning to make its transition to self-driving cars.

For this reason United States Department of Transportation is setting up central service to monitor traffic conditions nationwide

Sensors over the interstate system monitor traffic conditions like: car speeds, pavement and weather conditions, etc.

Interstate Highway System






- 16.000 sensors
- Measure
 - Speed
 - Travel Time
 - Weather, pavement, and traffic condition
- Support desktop, mobile, and car navigation systems




Other Requirements

- Need to keep 3 years of history
- Three data centres
 - New York
 - Chicago
 - Los Angeles
- Need to support 5 millions simultaneous users
 - Peak volume (rush hour)
 - Every minute, each request the 10 minute average speed for 50 sensors

A (discrete) *time series* is a sequence of data, measured typically at successive points in time, spaced at uniform time intervals.



[illegible]



Schema Design Consideration

- Horizontal Scalability
- Store event data
- Support analytical queries
- Find best compromise between:
 - Memory utilisation
 - Write performance
 - Read/Analytical query performance
 - Accomplish with a realistic amount of hardware



There are different choices in order to store upcoming information with a document DB:

- Document per event
- Document per minute (average)
- Document per minute (second)
- Document per hour
- ...



```
{  
  segID: "I80_mile34",  
  speed: 63,  
  ts: ISODate("2016-11-10T22:56:00.00-0500")  
}
```

- Relational centric approach
- Insert driven workload
- Overall performances and resources consumption are same as with a relational DB



Document per minute (average)

```
{  
  segID: "I80_mile34",  
  speed_num: 18,  
  speed_sum: 1256,  
  ts: ISODate("2016-11-10T22:56:00.00-0500")  
}
```

- Pre-aggregate to compute average per minute easily
- Update driven workload
- Resolution at the minute level



Document per minute (second)

```
{  
  segID: "I80_mile34",  
  speed: {0:63, 1:23, 2:45, ..., 59:65},  
  ts: ISODate("2016-11-10T22:56:00.00-0500")  
}
```

- Store per-second data at minute level
- Update driven workload
- Pre allocate structure to avoid document move



```
{  
  segID: "I80_mile34",  
  speed: {0:63, 1:23, 2:45, ..., 3599:65},  
  ts: ISODate("2016-11-10T22:00:00.00-0500")  
}
```

- Store per-second data at hourly level
- Update driven workload
- Pre allocate structure to avoid document move
- **Updating last second requires 3599 steps**



Document per hour (by second)

```
{  
  segID: "I80_mile34",  
  speed:{0: {0:63, ..., 59:45},  
        ...,  
        59:{0:65, ..., 59:65}  
  }  
  ts: ISODate("2016-11-10T22:00:00.00-0500")  
}
```

- Store per-second data at hourly level with nesting
- Update driven workload
- Pre allocate structure to avoid document move
- **Updating last second requires 59 + 59 steps**



Characterizing write differences

- Example: data generated every second
- For one minute:
 - *Document Per Event* \rightarrow 60 Writes
 - *Document Per Minute* \rightarrow 1 Write, 59 updates
- Transition from write driven to update driven
 - individual writes are smaller
 - performance and concurrency benefits



Characterizing read differences

- Example: data generated every second
- Reading data for a single hour requires:
 - *Document Per Event* \rightarrow 3600 reads
 - *Document Per Minute* \rightarrow 60 reads
- Read performance is greatly improved:
 - optimal with tuned block and read ahead
 - fewer disks seeks



Characterizing memory differences

- `_id` index for 1 billion events:
 - *Document Per Event* → 32 Gb
 - *Document Per Minute* → 0.5 Gb
- `_id` index plus `segId` and `ts` index:
 - *Document Per Event* → 100Gb
 - *Document Per Minute* → 2 Gb
- memory requirements significantly reduced:
 - fewer shards
 - lower capacity servers



- Writes:
 - 16.000 sensors, 1 update per minute
 - $16.000 / 60 = 267$ updates per second
- Reads:
 - 5 millions simultaneous users
 - Each request data for 50 sensors per minute

Query: Find the average speed over the last ten minutes

10 minute average query		
Schema	1 sensor	50 sensors
1 doc per event	10	500
1 doc per 10 min	1.9	95
1 doc per hour	1.3	65

10 minute average query with 5M users	
Schema	ops/sec
1 doc per event	42M
1 doc per 10 min	8M
1 doc per hour	5.4M

Writes: impacts of alternative schemas

1 Sensor - 1 Hour		
Schema	Inserts	Updates
doc/event	60	0
doc/10 min	6	54
doc/hour	1	59

16000 Sensors - 1 Day		
Schema	Inserts	Updates
doc/event	23M	0
doc/10 min	2.3M	21M
doc/hour	.38M	22.7M

Memory: impacts of alternative schemas

1 Sensor - 1 Hour		
Schema	# of Documents	Index Size (bytes)
doc/event	60	4200
doc/10 min	6	420
doc/hour	1	70

16000 Sensors – 1 Day		
Schema	# of Documents	Index Size
doc/event	23M	1.3 GB
doc/10 min	2.3M	131 MB
doc/hour	.38M	1.4 MB

```
{ _id: ObjectId("5382ccdd58db8b81730344e2"),  
  linkId: 900006,  
  date: ISODate("2014-03-12T17:00:00Z"),  
  data: [  
    { speed: NaN, time: NaN },  
    { speed: NaN, time: NaN },  
    { speed: NaN, time: NaN },  
    ...  
  ],  
  conditions: {  
    status: "Snow / Ice Conditions",  
    pavement: "Icy Spots",  
    weather: "Light Snow"  
  }  
}
```



Query: two indexes required

```
{
  "segId" : "20484097",
  "ts" : ISODate("2013-10-10T23:06:37.000Z"),
  "time" : "237",
  "speed" : "52",
  "pavement": "Wet Spots",
  "status" : "Wet Conditions",
  "weather" : "Light Rain"
}
```

~70 bytes per document



A Smart Sample Index

```
{ _id: "900006:14031217"  
  data: [  
    { speed: NaN, time: NaN },  
    { speed: NaN, time: NaN },  
    { speed: NaN, time: NaN },  
    ...  
  ],  
  conditions: {  
    status: "Snow / Ice Conditions",  
    pavement: "Icy Spots",  
    weather: "Light Snow"  
  }  
}
```

Saves an extra index



Sample index: range query

```
{ _id: "900006:14031217",  
  data: [  
    { speed: NaN, time: NaN },  
    { speed: NaN, time: NaN },  
    { speed: NaN, time: NaN },  
    ...  
  ],  
  conditions: {  
    status: "Snow / Ice Conditions",  
    pavement: "Icy Spots",  
    weather: "Light Snow"  
  }  
}
```

Range queries:
/^900006:1403/


Regex must be
left-anchored &
case-sensitive



Preallocate data

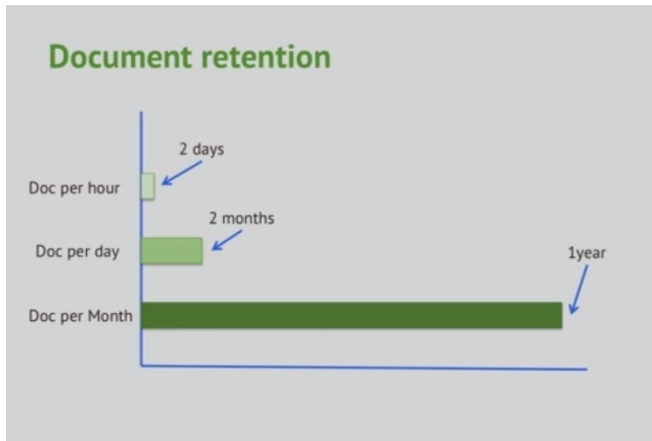
```
{ _id: "900006:140312",  
  data: [  
    { speed: NaN, time: NaN },  
    { speed: NaN, time: NaN },  
    { speed: NaN, time: NaN },  
    ...  
  ],  
  conditions: {  
    status: "Snow / Ice Conditions",  
    pavement: "Icy Spots",  
    weather: "Light Snow"  
  }  
}
```

Pre-allocated,
60 element array of
per-minute data





```
{ _id: "20484097:20140204",  
  hours: [  
    { speed: { sum: 1889, count: 60 }  
      time: { sum: 20562, count: 60 },  
      conditions: {  
        status: "Snow / Ice Conditions",  
        pavement: "Icy Spots",  
        weather: "Light Snow"  
      }  
    },  
    { speed: {m: 1892, count: 60 },  
      time: {sum: 20442, count: 60 },  
      conditions: {  
        status: "Snow / Ice Conditions",  
        pavement: "Slush",  
        weather: "Light Snow"  
      }  
    }  
  ]  
}
```





What is the average speed for a segment ?

```
>db.linkData.aggregate{
  {$match:{"_id":"/^20484097/"}} ,
  {$project:{"data.speed":1 ,linkId:1}} ,
  {$unwind:"$data"} ,
  {$group:{"_id":"$linkId",ave:$avg{"$data.speed"}}
};

{ "_id": "20484097", "ave": "47.3665646546" }
```