



Non è possibile visualizzare l'immagine.

Chapter 7: Entity-Relationship Model

E-R to Relational translation

Database System Concepts, 7th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use



Reduction to Relation Schemas



Reduction to Relation Schemas

- Entity sets and relationship sets can be expressed uniformly as *relation schemas* that represent the contents of the database.
- A database which conforms to an E-R diagram can be represented by a collection of schemas.
- For each entity set and relationship set there is a unique schema that is assigned the name of the corresponding entity set or relationship set.
- Each schema has a number of columns (generally corresponding to attributes), which have unique names.




Representing Entity Sets

- A strong entity set reduces to a schema with the same attributes

student(ID, name, tot_cred)

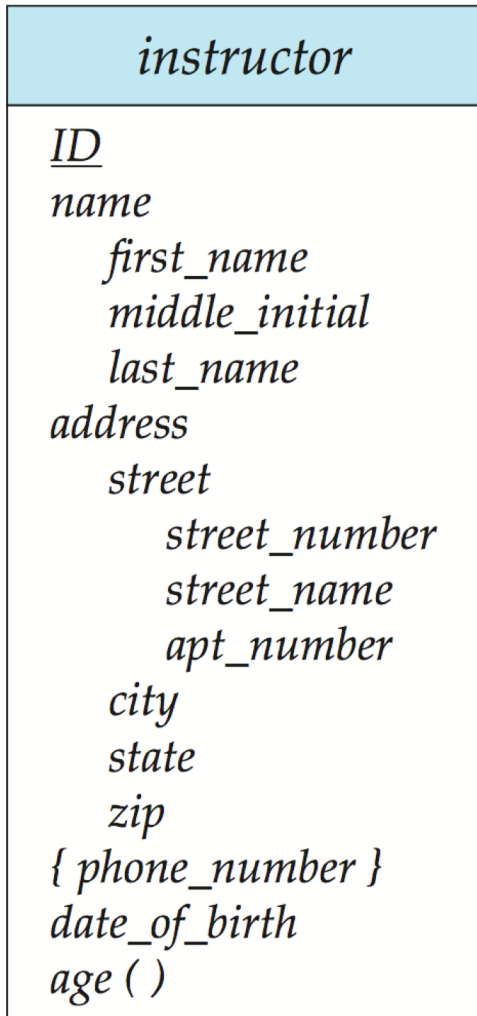
- A weak entity set becomes a table that includes a column for the primary key of the identifying strong entity set

section (course_id, sec_id, sem, year)





Representation of Entity Sets with Composite Attributes



- Composite attributes are flattened out by creating a separate attribute for each component attribute
 - Example: given entity set *instructor* with composite attribute *name* with component attributes *first_name* and *last_name* the schema corresponding to the entity set has two attributes *name_first_name* and *name_last_name*
 - ▶ Prefix omitted if there is no ambiguity (*name_first_name* could be *first_name*)
- Ignoring multivalued attributes, extended instructor schema is
 - *instructor*(ID, *first_name*, *middle_initial*, *last_name*, *street_number*, *street_name*, *apt_number*, *city*, *state*, *zip_code*, *date_of_birth*)



Representation of Entity Sets with Multivalued Attributes


- A multivalued attribute M of an entity E is represented by a separate schema EM
- Schema EM has attributes corresponding to the primary key of E and an attribute corresponding to multivalued attribute M
- Example: Multivalued attribute *phone_number* of *instructor* is represented by a schema:
$$inst_phone = (\underline{ID}, \underline{phone_number})$$
- Each value of the multivalued attribute maps to a separate tuple of the relation on schema EM
 - For example, an *instructor* entity with primary key 22222 and phone numbers 456-7890 and 123-4567 maps to two tuples:
(22222, 456-7890) and (22222, 123-4567)



Representing Relationship Sets

- A many-to-many relationship set is represented as a schema with attributes for the primary keys of the two participating entity sets, and any descriptive attributes of the relationship set.
- Example: schema for relationship set *advisor*

$advisor = (\underline{s_id}, \underline{i_id})$



A note on the ER-to-relational mapping: the case of one-to-one relationships


Angelo Montanari

Dept. of Mathematics, Computer Science, and Physics

University of Udine, Italy

Course on Data Management for Big Data

The mapping of one-to-one relationships



We have to distinguish among 3 different cases:

- ▶ $E1(\mathbf{PK1}, A1) - (1, 1) - R(A) - (0, 1) - E2(\mathbf{PK2}, A2)$

where $PK1$ is the key of entity $E1$, $A1$ is an attribute of $E1$, A is an attribute of relationship R , $PK2$ is the key of entity $E2$, and $A2$ is an attribute of $E2$

(the case

$$E1(\mathbf{PK1}, A1) - (0, 1) - R(A) - (1, 1) - E2(\mathbf{PK2}, A2)$$

is completely symmetric, and thus ignored)

- ▶ $E1(\mathbf{PK1}, A1) - (0, 1) - R(A) - (0, 1) - E2(\mathbf{PK2}, A2)$
- ▶ $E1(\mathbf{PK1}, A1) - (1, 1) - R(A) - (1, 1) - E2(\mathbf{PK2}, A2)$

The case of $(1, 1) - (0, 1)$ relationships

How can we map the following ER schema into a corresponding relational one (introducing no redundancy, and preserving as much as possible information/constraints)?

$E1(\mathbf{PK1}, A1) - (1, 1) - R(A) - (0, 1) - E2(\mathbf{PK2}, A2)$

The case of $(1, 1) - (0, 1)$ relationships

How can we map the following ER schema into a corresponding relational one (introducing no redundancy, and preserving as much as possible information/constraints)?

$E1(\mathbf{PK1}, A1) - (1, 1) - R(A) - (0, 1) - E2(\mathbf{PK2}, A2)$

Relational schema:

$E1(\underline{PK1}, A1, PK2, A)$ and $E2(\underline{PK2}, A2)$, where $PK2$ is a foreign key of relation $E1$ that refers to the primary key $PK2$ of relation $E2$.

The case of $(1, 1) - (0, 1)$ relationships

How can we map the following ER schema into a corresponding relational one (introducing no redundancy, and preserving as much as possible information/constraints)?

$E1(\mathbf{PK1}, A1) - (1, 1) - R(A) - (0, 1) - E2(\mathbf{PK2}, A2)$

Relational schema:

$E1(\underline{PK1}, A1, PK2, A)$ and $E2(\underline{PK2}, A2)$, where $PK2$ is a foreign key of relation $E1$ that refers to the primary key $PK2$ of relation $E2$.

- ▶ $(1, _)$ on $E1$ side: $PK2$ NOT NULL in $E1$
- ▶ $(_, 1)$ on $E1$ side: $PK1$ is the primary key
- ▶ $(0, _)$ on $E2$ side: $PK2$ foreign key in $E1$ referring to the primary key of $E2$
- ▶ $(_, 1)$ on $E2$ side: $PK2$ UNIQUE in $E1$

The case of $(0, 1) - (0, 1)$ relationships

$$E1(\mathbf{PK1}, A1) - (0, 1) - R(A) - (0, 1) - E2(\mathbf{PK2}, A2)$$

The case of $(0, 1) - (0, 1)$ relationships

$$E1(\mathbf{PK1}, A1) - (0, 1) - R(A) - (0, 1) - E2(\mathbf{PK2}, A2)$$

Relational schema:

$E1(\underline{PK1}, A1, PK2, A)$ and $E2(\underline{PK2}, A2)$, where $PK2$ is a foreign key of relation $E1$ that refers to the primary key $PK2$ of relation $E2$.

The case of $(0, 1) - (0, 1)$ relationships

$E1(\mathbf{PK1}, A1) - (0, 1) - R(A) - (0, 1) - E2(\mathbf{PK2}, A2)$

Relational schema:

$E1(\underline{PK1}, A1, PK2, A)$ and $E2(\underline{PK2}, A2)$, where $PK2$ is a foreign key of relation $E1$ that refers to the primary key $PK2$ of relation $E2$.

- ▶ $(0, _)$ on $E1$ side: $PK2$ **can be NULL** in $E1$
- ▶ $(_, 1)$ on $E1$ side: $PK1$ is the primary key
- ▶ $(0, _)$ on $E2$ side: $PK2$ foreign key in $E1$ referring to the primary key of $E2$
- ▶ $(_, 1)$ on $E2$ side: $PK2$ UNIQUE in $E1$

$E1(\underline{PK1}, A1)$ and $E2(\underline{PK2}, A2, PK1, A)$, where $PK1$ is a foreign key of relation $E2$ that refers to the primary key $PK1$ of relation $E1$, works as well.

The case of $(0, 1) - (0, 1)$ relationships (contr'n'd)

How do we choose between the two options? Participation of entities in the relationship can be a criterion.

In case the participation of both entities in the relationship is very low, a third option is possible

The case of $(0, 1) - (0, 1)$ relationships (contr'n'd)

How do we choose between the two options? Participation of entities in the relationship can be a criterion.

In case the participation of both entities in the relationship is very low, a third option is possible

Relational schema:

$E1(\underline{PK1}, A1)$, $E2(\underline{PK2}, A2)$, and $R(\underline{PK1}, PK2, A)$ (or $R(PK1, \underline{PK2}, A)$), where $PK1$ is a foreign key of relation R that refers to the primary key $PK1$ of relation $E1$ and $PK2$ is a foreign key of relation R that refers to the primary key $PK2$ of relation $E2$.

The case of $(0, 1) - (0, 1)$ relationships (contr'n'd)

How do we choose between the two options? Participation of entities in the relationship can be a criterion.

In case the participation of both entities in the relationship is very low, a third option is possible

Relational schema:

$E1(\underline{PK1}, A1)$, $E2(\underline{PK2}, A2)$, and $R(\underline{PK1}, PK2, A)$ (or $R(PK1, \underline{PK2}, A)$), where $PK1$ is a foreign key of relation R that refers to the primary key $PK1$ of relation $E1$ and $PK2$ is a foreign key of relation R that refers to the primary key $PK2$ of relation $E2$.

- ▶ $(0, _)$ on $E1$ side: $PK1$ foreign key in R referring to the primary key of $E1$
- ▶ $(_, 1)$ on $E1$ side: $PK1$ is the primary key of R
- ▶ $(0, _)$ on $E2$ side: $PK2$ foreign key in R referring to the primary key of $E2$
- ▶ $(_, 1)$ on $E2$ side: $PK2$ UNIQUE in R

The case of $(1, 1) - (1, 1)$ relationships

No one of the previous solutions works with

$$E1(\mathbf{PK1}, A1) - (1, 1) - R(A) - (1, 1) - E2(\mathbf{PK2}, A2)$$

The case of $(1, 1) - (1, 1)$ relationships

No one of the previous solutions works with

$E1(\mathbf{PK1}, A1) - (1, 1) - R(A) - (1, 1) - E2(\mathbf{PK2}, A2)$

Relational schema:

$R(\underline{PK1}, A1, PK2, A2, A)$ (or $R(PK1, A1, \underline{PK2}, A2, A)$)

The case of $(1, 1) - (1, 1)$ relationships

No one of the previous solutions works with

$E1(\mathbf{PK1}, A1) - (1, 1) - R(A) - (1, 1) - E2(\mathbf{PK2}, A2)$

Relational schema:

$R(\underline{PK1}, A1, PK2, A2, A)$ (or $R(PK1, A1, \underline{PK2}, A2, A)$)

- ▶ $(1, _)$ on $E1$ side: $PK2$ NOT NULL in R
- ▶ $(_, 1)$ on $E1$ side: $PK1$ is the primary key (UNIQUE)
- ▶ $(1, _)$ on $E2$ side: $PK1$ is the primary key (NOT NULL)
- ▶ $(_, 1)$ on $E2$ side: $PK2$ UNIQUE in R

The case of one-to-many and many-to-many relationships

A similar analysis can be done for the cases of both one-to-many and many-to-many relationships

Such an analysis allows one to determine

- ▶ which constraints can be directly encoded in the relational schema, and
- ▶ which ones need to be explicitly forced.



Redundancy of Schemas

- Many-to-one and one-to-many relationship sets that are total on the many-side can be represented by adding an extra attribute to the “many” side, containing the primary key of the “one” side
- Example: Instead of creating a schema for relationship set *inst_dept*, add an attribute *dept_name* to the schema arising from entity set *instructor*

