



DMIF, University of Udine

DMBD – Part I

Introduction

Andrea Brunello

andrea.brunello@uniud.it



- The need for databases
 - *Why can't we simply store files on the disk?*
- Abstraction levels and data models
 - *Things that allow us to define and organize data*
- Database design
 - *The steps to follow to build a database*
- Relational databases
 - *The evergreen way-to-go for databases*
- Database internals
 - *How is data stored? How can we interact with a database?*
- History of database systems
 - *When did it all start?*



Database Management System (DBMS)

- A DMBS contains information on a particular domain
 - Collection of interrelated data (database)
 - Set of programs to manage the data:
 - Definition of structures to store data
 - Mechanisms to retrieve/modify data
 - Safety and concurrency aspects
- Database applications:
 - Banking: bank transfers, interactions with the ATMs
 - Airlines: reservations, schedules
 - Universities: student registrations, grade assignment
 - Sales: customers, inventory, products and sales
 - E-commerce: e.g, when you buy something from Amazon
- In general, databases can be very large
- They touch all aspects of our lives, although user interfaces hide access details



Example: University database

- A University database could support interactions such as:
 - Add new students, instructors, and courses
 - Register students for courses, and generate class rosters
 - Plan exams, assign grades to students, compute grade point averages
- In the early days, such applications were built directly on top of the file system
- Ad-hoc solutions: data stored in files (e.g., .csv) and applications coded in scripts
 - Over time, they both tend to grow in their number
- First relational database in early 80s, after 10 years of work
 - So, the need for a DBMS was seen from the beginning



Drawbacks of using file system to store data

- Data redundancy
 - Multiple file formats and programming languages
 - Duplication of information in different files
 - Other than wasting memory, brings to inconsistencies
- Difficulty in accessing data
 - Need to write a new program to carry out each new task
 - E.g., a program to get info about all professors of Computer Science, another for those teaching Economy, ...
 - Also, data access is not “interactive”
- Data isolation
 - Multiple files and formats are difficult to manage
- Constraints management
 - Integrity constraints (e.g., account balance > 0) become “buried” in program code rather being stated explicitly
 - Hard to add new constraints or change existing ones



Drawbacks of using file system to store data (Cont.)

- Atomicity of updates
 - Failures may leave the database in an inconsistent state with partial updates carried out
 - E.g., transfer of funds from one account to another should either complete or not happen at all
- Concurrent access by multiple users
 - What happens if two users try to access and modify the same data?
 - E.g., two users booking the last airline ticket
 - For efficiency reasons, we must support concurrency
- Security problems
 - Hard to provide a user access to some, but not all, data

⇒ DBMSs provide solutions to all these problems



Characteristics justifying the use of a DBMS

- Large amounts of data
 - Not possible to load all of them into *main memory*
 - Critical complexity factor for a DBMS is the number of transfers to/from disk
- Global data
 - DB and its data are relevant for a vast array of users and applications
 - They are typically not tied to a particular user or application, as standard programs are
- Persistence of data
 - DB is there independently from the interacting users or applications
 - Data lives on its own unlike, for instance, program variables




- To tackle the complexity inherent in their development, DBMSs make use of **abstractions**
- Abstractions are a fundamental concept in computer science, for instance, in the Operating Systems domain
- They keep complexity under control allowing the user to focus on what is important, leaving out irrelevant details that would generate confusion
- In practice, for what concerns the database realm, it involves stratifying the DBMS



DBMS levels of abstraction

- From lowest to highest:
 - **Physical level:** describes how data is stored on the disk
 - **Logical level:** describes how data is stored and represented in the database and their relationships (data model)
 - E.g., by means of tables in the relational model
 - They are *tables*, we do not care about how they are stored on the disk (the physical level deals with it)
 - **View/External level:** manages the presentation of information to users and programs. Views can also hide information to some kinds of users
- Each level encapsulates/hides the details it deals with
- Each level can take advantage of the functionalities made available by the lower levels, without caring about how they are implemented



- **Physical data independence:** ability to make changes to the physical level without having to modify the logical one
- It's a concept similar to the APIs (Application Programming Interface)



Instances and schemas

- The levels of abstraction of a database are described by means of schemas:
 - **Physical schema:** it lays out how data are stored physically on a storage system in terms of files and indices
 - **Logical schema:** it encodes, by means of a *data model*, the overall logical structure of the database
 - E.g., the database consists of information about a set of customers (each characterized by name, address, and phone number) and their accounts in a bank; the relationships among them; and, some constraints over them
 - It allows to describe at an abstract level the characteristics of the elements I am interested in dealing with
- While the schemas lay down the “structure of the database” at different levels, the **instance** of the database refers to its actual content at a particular moment
 - E.g., the specific, concrete customers and accounts



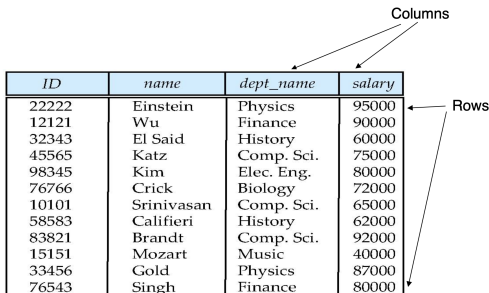
- A collection of tools for describing:
 - Data (and its structure)
 - Data relationships
 - Data constraints
- There are two main kinds of data models:
 - **Conceptual data models:** they describe in an abstract way a domain without reference to a specific “database technology” (Entity-Relationship model, for DB design)
 - **Logical data models:** they are specific of a given database technology (although still independent from the physical details)
 - Relational model
 - Graph data model
 - Object-based data models
 - Semistructured data models (XML)
 - Older models: Network model, Hierarchical model

The process of designing a database is typically articulated into different phases:

- ① Brainstorming meetings with IT personnel and all interested stakeholders
 - Collection of requirements
 - Design of a conceptual model (e.g., E-R model)
 - The E-R model has a specific notation, the *E-R diagram*, that helps all involved parts to discuss about the future database
- ② Translation of the conceptual model into a logical model
 - Typically, a set of translation rules is followed
 - At this stage, a specific DBMS technology must be chosen (e.g., relational DB)
- ③ Addition to the logical model of details regarding the physical, low-level aspects (e.g., usage of indexes)
 - The physical schema is thus obtained

Relational model

- In the relational model, all data is stored by means of tables
- Such tables represent both data and the relationships among data
- It is a kind of **record-based** model: rows have a fixed format and length
- Each table has a **primary key**



<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

(a) The *instructor* table



Relational model

Encoding of relationships

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

(a) The *instructor* table

<i>dept_name</i>	<i>building</i>	<i>budget</i>
Comp. Sci.	Taylor	100000
Biology	Watson	90000
Elec. Eng.	Taylor	85000
Music	Packard	80000
Finance	Painter	120000
History	Painter	50000
Physics	Watson	70000

(b) The *department* table



Another example

What is wrong with this table?

<i>ID</i>	<i>name</i>	<i>salary</i>	<i>dept_name</i>	<i>building</i>	<i>budget</i>
22222	Einstein	95000	Physics	Watson	70000
12121	Wu	90000	Finance	Painter	120000
32343	El Said	60000	History	Painter	50000
45565	Katz	75000	Comp. Sci.	Taylor	100000
98345	Kim	80000	Elec. Eng.	Taylor	85000
76766	Crick	72000	Biology	Watson	90000
10101	Srinivasan	65000	Comp. Sci.	Taylor	100000
58583	Califieri	62000	History	Painter	50000
83821	Brandt	92000	Comp. Sci	Taylor	100000
15151	Mozart	40000	Music	Packard	80000
33456	Gold	87000	Physics	Watson	70000
76543	Singh	80000	Finance	Painter	120000



Another example

What is wrong with this table? (Cont.)

- There are row insert/update/delete anomalies, for instance:
 - Cannot insert an employee without a linked project
 - Cannot keep a project if it does not have any employees
 - Potential need to modify multiple rows to update the salary of an employee or the budget of a project
- Intuitively, the problem is that the table does not have a clear semantics
 - It is mixing information about employees and departments
 - The typical solution involves splitting the table
- **Normalization theory** provides formal techniques to prevent this kind of problems




Structured Query Language

- A very nice thing about the relational DBs is that they come with SQL (Structured Query Language)
- SQL is a *declarative* language composed of two parts:
 - DDL (Data Definition Language): that allows to specify the schema of the database
 - DML (Data Manipulation Language): that allows to interact with the content of the database (i.e., the database instance) by means of *queries*
- SQL features a very intuitive syntax
- It is standardized within ISO/IEC, so to allow for a better interoperability between different relational DBMSs

```
SELECT name  
FROM instructor  
WHERE salary > 60000;
```

Database system internals

Graphical account





- Other than data itself, the database contains several other information:
 - **Indices:** auxiliary structures that speed up data access when certain conditions are satisfied
 - **Data dictionary:** metadata that define the structure of the database
 - **Statistical data:** regarding the content of the database, the distribution of data, etc.

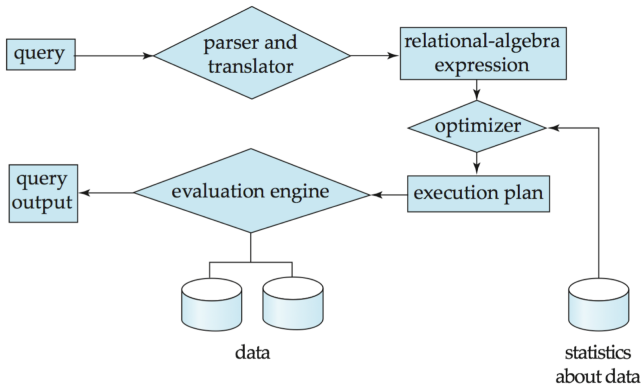


- From the point of view of its implementation, the DBMS is composed of a set of sub-modules, each dealing with a specific aspect of data and their access
 - Storage manager
 - Query processor
 - Transaction manager



- The storage manager is a program module that provides the interface between the low-level data stored in the database and the application programs and queries submitted to the system
- It is in charge of the following tasks:
 - Interaction with the OS file manager
 - Efficient storage, retrieval, and update of data
 - In the case of SQL, translation of DML statements into low-level file-system commands
- Deals with:
 - Storage access
 - File organization
 - Indexing

- Steps:
 - 1 Parsing and translation
 - 2 Optimization
 - 3 Evaluation





Query processing (Cont.)

- In general, there are many alternative ways of evaluating a given query
 - Equivalent expressions: e.g., first select the row with *ID* = 22222 from *instructor*, then remove all columns except for *ID* and *name*, or vice-versa
 - Different algorithms for each operation: e.g., perform a table scan or rely on indexes
- The cost difference (in terms of execution time and disk transfers) between a good and a bad way of evaluating a query can be enormous
- Need to estimate the cost of operations:
 - Heavily depends on statistical information about relations which the database must maintain (e.g., number of rows in a table)
 - Need to estimate statistics for intermediate results to compute the cost of complex expressions



Transaction manager

- What if the system fails?
- What if more than one user is concurrently updating the same data?
- A **transaction** is a collection of operations that performs a single logical function in a database application
 - E.g., a money transfer between two bank accounts
- The transaction manager ensures that the DB remains in a consistent state despite of system failures (e.g., power failures or OS crashes) and transaction failures
- In addition, interactions among concurrent transactions is monitored, to ensure the consistency of the database
 - We do not want transactions to interfere with one another in undesired ways




- **Database administrator:**
 - defines/manages the DB schema
 - manages data access
 - specifies integrity constraints
- **Database users:**
 - *basic users*: interact with the DB through user interfaces
 - *programmers*: access the DB when writing applications using high-level languages (e.g., Java, Python)
 - *advanced users*: directly query the DB through DML



- The architecture of a database system is greatly influenced by the underlying computer system on which the database is running:
 - Centralized
 - Client-server
 - Parallel (multi-processor)
 - Distributed
- In this part of the course we will not deal with these aspects, which are included in the second part

- 1950s and early 1960s:
 - Data processing done using magnetic tapes for storage, which provided only *sequential access*
 - Punched cards for input
 - Support for tasks such as payroll management





History of database systems (Cont.)

- Late 1960s and 1970s:
 - Hard disks allow for a *direct access* to data
 - Network and hierarchical data models become of widespread use
 - They use graphs and trees as data structures to store data
 - Ted Codd defines the relational data model
 - Would win the ACM Turing Award for this work
 - Points of strength: simplicity and possibility of hiding physical details
 - IBM Research begins work on System R prototype
 - UC Berkeley begins work on Ingres prototype
 - They both still exhibit computationally worse performance with respect to network and hierarchical data models
 - OLTP: On-Line Transaction Processing



History of database systems (Cont.)

- 1980s:
 - Relational DB prototypes evolve into commercial systems
 - SQL becomes industrial standard: programmers can now work at the logical level, without worrying about query optimality and the actual location of data
 - Parallel and distributed database systems
 - Object-oriented database systems: they use exactly the same model as object-oriented programming languages
- 1990s:
 - Data warehouses: subject oriented, integrated, time-variant, and non-volatile collection of data
 - Decision support systems and data mining applications
 - OLAP: On-Line Analytical Processing



History of database systems (Cont.)

- 2000s:
 - Structured, semi-structured, and unstructured data
 - XML and XQuery: handling of semi-structured data
 - Automated database administration
- Today:
 - Giant data storage systems (e.g., Google BigTable)
 - Hadoop framework to handle all aspects of data management



- MySQL: <https://www.mysql.com/it/>
- Oracle: <https://www.oracle.com/database/technologies/>
- PostgreSQL: <https://www.postgresql.org/>
- Microsoft SQL Server: <https://www.microsoft.com/en-us/sql-server/sql-server-downloads>