



DMIF, University of Udine

---


# An Introduction to Neural Networks

Andrea Brunello, Ph.D. Student  
andrea.brunello@uniud.it

November 22, 2018

A brain neuron can be seen as an organic switch:

- it receives multiple signals through *synapses*, located on the *dendrites*
- when the received signals are strong enough, the neuron is *activated* and emits a signal through the *axon*, which might then activate other neurons




An artificial neuron can be seen as an *abstraction* of a brain neuron:

- it computes a *weighted sum* of the inputs:

$$b + w_1x_1 + w_2x_2 + \cdots + w_nx_n$$

- the result is then passed through a nonlinear *activation function*



# Neural network, single output

The neurons are organized in layers ( $\leadsto$  abstraction hierarchies).

**Input:**  $x_1, \dots, x_n$  attribute values of the specific instance





Figure: A 2-layer, fully-connected, feed-forward neural network.

# Neural network, multiple outputs


**Multi-Output Neural Network**





# Activation function


- In classification tasks, to simulate a biological neuron, the activation func. should have a “switch on” characteristic
- *Sigmoid* was commonly used:  $S(x) = 1/(1 + e^{-x})$
- In this case, a single neuron is similar to a logistic classifier
- *Vanishing gradient* problem: the function saturates, i.e., in some regions, even a large change in the input will produce a small change in the output (gradient is small)



The choice of the most appropriate activation function depends on the specific task at hand


## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




## tanh

$$\tanh(x)$$




## ReLU

$$\max(0, x)$$



## Leaky ReLU

$$\max(0.1x, x)$$




## Maxout (Layer)

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

## ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$

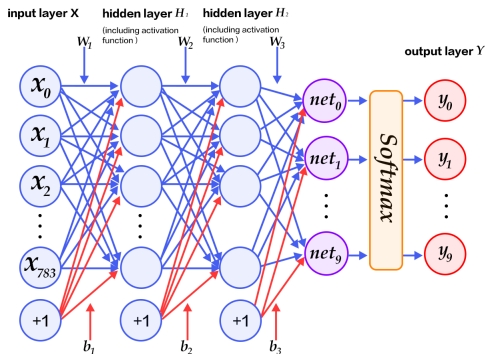


Typically, layers are homogeneous w.r.t the activation function

# Output layer and activation function

For the output layer, typically:

- **identity**: regression tasks
- **sigmoid**: multiple classification, e.g., image object recognition
- **softmax (layer)**: classification, it generates probabilities







# Properties of activation functions

Among the others:

- **continuously differentiable:** to allow for the training of the network (ReLU?)
- **nonlinear:** then, a two-layer (1 HL) neural network can be proven to be a *universal function approximator*: a FF network with 1 HL can represent any continuous function that maps intervals of  $\mathbb{R}$  to intervals of  $\mathbb{R}$ , but the layer may be infeasibly large and may fail to learn and generalize
- **monotonic:** in such case, the error surface associated with a single-layer model is guaranteed to be convex

Full list: [https://en.wikipedia.org/wiki/Activation\\_function](https://en.wikipedia.org/wiki/Activation_function)



- The weights and the biases in the network are initialized with *small random values*
- Then, they are *iteratively adjusted* considering the output of the net, through the following training loop:
  - 1 draw a batch of training samples  $x$  with corresponding targets  $y$
  - 2 run the network on  $x$  (forward pass) to obtain  $\hat{y}$
  - 3 compute a measure of the mismatch between  $y$  and  $\hat{y}$
  - 4 update all weights in the net in a way that slightly reduces the loss on this batch



# Training of the net: problem definition

- We want to find the best weights  $w$  and biases  $b$  so that the output  $a_i$  from the network approximates the label  $y(x_i)$  for every input  $x_1, x_2, \dots, x_n$
- To quantify how well we are achieving this goal, we define a *cost function*:

$$C(\vec{w}, \vec{b}) = \frac{1}{2n} \sum_{i=1}^n (y(x_i) - a_i)^2$$

- The objective is to minimize such cost function (in the biggest neural networks it may depend on billions of variables!)



# Training of the net: a simple strategy

How can we adjust the weights based on the result of the cost function on a batch?

- freeze all weights except for the one being considered, say  $w_k = 0.5$
- calculate the cost function for  $w_k = 0.75$  and  $w_k = 0.25$
- update the value of  $w$  setting it to the one that reduces the cost most
- repeat for all coefficients in the network (thousands or millions)


↪ Horribly inefficient!

Fortunately, we can take advantage of the fact that all operations used in the network are *differentiable*

# Training of the net: gradient descent

*Gradient descent* is used to determine how to vary the weights and biases, in order to minimize the cost

- the gradient (i.e., derivative) of the cost function w.r.t. the weights and biases is considered to find a (local) minimum
- single-weight case:  $w_{new} = w_{old} - \eta * \nabla C$ ,  $\nabla C = \frac{\partial C}{\partial w}$



In general, the gradient is a vector of partial derivatives. For each weight  $w^i$  and bias  $b^j$  of the net, we have:


$$w_{new}^i = w_{old}^i - \eta \frac{\partial C}{\partial w_{old}^i}$$
$$b_{new}^j = b_{old}^j - \eta \frac{\partial C}{\partial b_{old}^j}$$

How may we calculate the gradient of the cost function, i.e., the partial derivatives with respect to any weight and bias in the network?

- *Backpropagation* algorithm (chainning rule)

As all other machine learning techniques, also neural networks are prone to overfitting:

- it happens when the model is too complex w.r.t. the amount of training data available
- the performance of the model on validation data peaks after a few epochs, then begins to degrade





# Overfitting: possible solutions

How to counter overfitting?

- increase the amount of training data
- decrease the model complexity (number and size of layers)
- reduce the number of training iterations (epochs)
- specific strategies:
  - *regularization*
  - *dropout*





# Overfitting: regularization

Idea: a *simple mode* is a model where the distribution of parameter values is “smooth”:

- a possible way to achieve this is by forcing the weights to take only small values
- this typically makes the distribution of weight values more regular (weight regularization)
- simply add to the loss function of the network a cost associated to having large weights
- two main flavours:
  - *L1\_regularization*: the cost added is proportional to the absolute value of the weight coefficients
  - *L2\_regularization*: the cost added is proportional to the square of the value of the weight coefficients




# Overfitting: dropout

One of the most effective and commonly used regularization techniques for neural networks:


- during *training*, a fraction of output features of the layer (typically, 0.2-0.5) are randomly dropped out (set to zero)
- at *test* time, no units are dropped. Instead, the layer's output values are scaled down by a factor equal to the dropout rate, to balance for the fact that more units are active than at training time
- the core idea is that introducing noise in the output values of a layer can break up patterns that are not significant, which the network would start memorizing if no noise is present

- classical, fully-connected layers learn *global patterns* in their input space
- sometimes, it is useful to learn *local patterns*
- for example, when dealing with image classification, we may extract patterns referred to small 2D patches of the input image
- typical operations: *convolution* + *pooling* (downsampling)


7 x 7 Input Volume



5 x 5 Output Volume



Single depth slice




max pool with 2x2 filters  
and stride 2



Note: there are also 1D convs, typically applied to seq. data

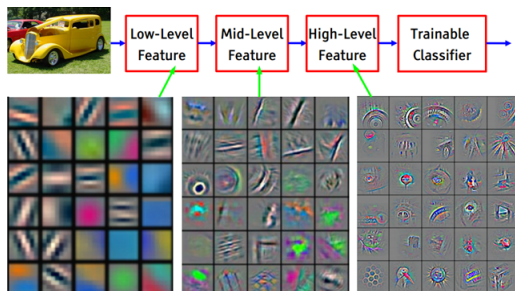
# Convolutional Neural Networks

- 2D convolutions typically operate over 3D data (height, width, color depth of an image)
- the output is still 3D, with an arbitrary depth
- intuitively, each output channel in the depth axis works as a *filter*, encoding specific aspects of the input data
- the convolutional layers are then followed by a series of classic, fully-connected layers



CNNs have two interesting properties:

- the patterns they learn are *translation invariant*: a pattern learnt in the lower-right corner of a picture can be recognized everywhere
- they learn *spatial hierarchies* of patterns






All neural networks described so far *have no memory* (FF):

- with them, in order to process a sequence or a temporal series of data points, you have to show the entire sequence to the network at once, i.e., turn it into a single data point

A RNN processes sequences by iterating through the sequence elements:

- it keeps a *state* containing information on what it has seen so far in a sequence
- the state is reset between processing two different, independent sequences

In principle, a RNN consists of a *for loop* that reuses quantities computed during the previous iteration of the loop




Although it should theoretically be able to retain at time  $t$  information about inputs seen many timesteps before, in practice such long-term dependencies are impossible to learn, due to the *vanishing gradient problem*

A solution to such a problem is given by specifically designed units, such as:

- Long Short-Term Memory (LSTM)
- Gated Recurrent Unit (GRU)

In short, a LSTM units allows the net to save information for later use, thus preventing older signals from gradually vanishing during processing







# Some history: the first models

1943 Neurophysiologist S. McCulloch and mathematician W. Pitts develop the model of the first **artificial neuron**, subsequently implemented by a simple circuit

- input could be 0 or 1, and excitatory (+) or inhibitory (-)
- inputs are summed considering their sign
- if the sum surpassed a certain threshold, output was 1, otherwise it was set to 0
- no learning mechanism



1949 D. Hebb proposes **connectionism** theory

- neural pathways strengthen over each successive use,
- especially between neurons that tend to fire at the same time

# Some history: the Perceptron

1957 F. Rosenblatt develops the **Perceptron**, which is the first “practical” artificial neural network

- single-layer network (in fact, a single neuron)
- linear classifier (e.g., can calculate AND and OR)
- weights are self-adjusted during the training phase
- “The Navy revealed the embryo of an electronic computer today that it expects will be able to walk, talk, see, write, reproduce itself and be conscious of its existence [...] Dr. F. Rosenblatt said Perceptrons might be fired to the planets as mechanical space explorers” (NYT)





# Some history: the Golden Age

1959 At Stanford, B. Widrow and M. Hoff develop **ADALINE**

- single-layer artificial neural network with multiple nodes
- each node accepts multiple inputs and generates an output
- improved training approach w.r.t. the Perceptron, that takes into account how much the error changes when each weight is changed (i.e., the derivative)

1962 **MADALINE** is the first neural network successfully applied to a real world problem

- two-layer, fully connected artificial neural network of ADALINE units
- removal of noise in phone lines



# Some history: the AI Winter

1969 M. Minsky (founder of the MIT AI Lab) and S. Papert (director of the lab) publish the book **Perceptrons**

- they prove that basic Perceptrons are incapable of processing the XOR circuit
- they further argue that the Perceptron could not be translated into multi-layered neural networks

1970s The **AI Winter**

- neural network research slowdown
- public and private institutions abandon neural network approaches, focusing instead on *expert systems*
- in 1975, P. Werbos discovers the potential of *Backpropagation* in neural networks, and goes unnoticed



# Some history: the rebirth

- 1980 K. Fukushima proposes the **Neocognitron**, which served as the inspiration for convolutional neural networks
- 1982 **Hopfield network** is one of the first forms of RNN
- 1986 D. Rumelhart, G. Hinton, and R. Williams publish **Learning representations by back-propagating errors**
- they clearly and concisely state the idea of Backpropagation, which finally becomes widely known
  - also, they specifically address the problems discussed by Minsky in Perceptrons
- 1989 It is mathematically proven that multiple layers allow neural networks to act as **universal approximators**
- 1989 AT&T Bell Labs publish the work **Backpropagation Applied to Handwritten Zip Code Recognition**



## Some history: some problems emerge

- 1991 Hochreiter identifies the **vanishing and exploding gradient** problems of backpropagation, that make it difficult to train f.f. deep or recurrent neural networks
- 1995 S. Thrun investigates the problem of learning a neural network capable of playing chess, getting far worse results than those of the standard GNU-Chess software: **computers are still too weak**
- 1995 LeCun et al. find that **Support Vector Machines** worked better or the same as all but the best designed neural nets, for the task of handwritten digit recognition



## Some history: another setback

1997 Schmidhuber and Hochreiter introduce **Long Short Term Memory** (LSTM), that essentially solves the problem of how to train recurrent neural networks

Late 90s Neural networks are still seen as a hassle to work with - the computers were not fast enough, and the algorithms were not smart enough

2001 **Random Forests** show remarkable performances, and are easier to work with than neural networks, which are once again disavowed by the machine learning community



# Some history: the second rebirth

**2006** Hinton, Osindero and Teh publish **A fast learning algorithm for deep belief nets**

- neural nets are rebranded as **Deep Learning**
- a new way of initializing the weights is presented, that allows to train neural networks with many layers well
- the key is having many layers of computing units so that good high-level representation of data could be learned, in contrast to hand-designing them

**2009** Hinton et al. develop a neural network capable of superseding the existent models for phoneme recognition


**2009** GPU-powered training

**2012** **Microsoft, Google, IBM** and Hinton's lab publish "Deep Neural Networks for Acoustic Modeling in Speech Recognition: The Shared Views of Four Research Groups"



# Some history: the final crowning

In **2012**, a convolutional neural network performs way better than previous approaches for image classification





# The end (?)

So, why did purely supervised learning with backpropagation not work well in the past? Geoffrey Hinton summarized the findings up to today in these four points:

- ① Our labeled datasets were thousands of times too small.
- ② Our computers were millions of times too slow.
- ③ We initialized the weights in a stupid way.
- ④ We used the wrong type of non-linearity.



... well, probably not. What about the future?

Deep learning models have some *severe limitations*:

- the only real success so far has been the ability to map a space  $X$  to space  $Y$  using a continuous geometric transform, given huge amounts of annotated data
- i.e., there isn't any kind of reasoning, nor understanding behind: it is just a chain of simple, continuous geometric transformations mapping one vector space into another
- a deep learning model can be interpreted as a kind of program, but, inversely, most programs cannot be expressed as deep learning models
- e.g., even learning a sorting algorithm with a deep neural network is tremendously difficult



In summary:

- deep learning models are severely limited in what they can represent
- given a task, even if an equivalent deep learning model does exist, it may be far too complex or simply not learnable, due to the current learning algorithms or lack of proper data
- we are still a long way from human-level AI

⇒ however, there's some light ahead...



Consider RNNs:

- they have slightly fewer limitations than FF networks, because they are a bit more than mere geometric transformations
- they are geometric transformations *repeatedly applied inside a for loop*, though a hardcoded one, and with severe limitations on what it does

We can build on such an idea:

- imagine a neural network augmented with programming primitives
- general for loops, if branches, while statements, variable creation, disk storage for long-term memory (NTUs), sorting operations, advanced data structures, ...



# Models as programs

The space of programs that such a network could represent would be far broader, however:

- the model would be mostly no longer differentiable, even though some parts of it may still be (subroutines)
- new, complementary training strategies will be required, for example based on *evolutionary algorithms*

Once learned, subroutines may be stored in a kind of repository, so to reuse them when needed for another task, achieving higher generalization and abstraction than pretrained weights

▷ Contemporary RNNs can be seen as a prehistoric ancestor of such hybrid algorithmic-geometric models, that can be interpreted as an *artificial general intelligence* (AGI)