



University of Udine

---

# Monitors that Learn from Failures

*Machine Learning-based Monitoring  
for Runtime System Verification*

Andrea Brunello - [andrea.brunello@uniud.it](mailto:andrea.brunello@uniud.it)




In several domains, systems generate **continuous streams of data** which may contain useful telemetry information

- They can be used for tasks such as predictive maintenance and preemptive failure detection (Industry 4.0)
- System behaviours can be convoluted, being the result of the interaction among several components and the environment
- Given the complexity of this setting, **deep learning** approaches are typically been considered. Problems:
  - resulting models are hardly interpretable
  - difficulty in providing guarantees on the obtained results

In critical contexts, formal methods have been recognized as an effective approach to ensure the correct behaviour of a system.

However, classical techniques, such as model checking, require a **complete specification** of the system and of the properties to be checked against it, and work in an **offline** fashion.

-> In some cases, their application can be very difficult!






Framework that **combines machine learning and monitoring** to detect critical system behaviours in an on-line setting:

- system behaviour's complexity is dealt with by means of machine learning
- extracted formal properties are interpretable, so a domain expert can easily read and validate the generated model
- the framework is highly modular with respect to the logic used to encode the system properties

Monitoring is a **run-time** verification technique:

- it establishes satisfaction/violation of a property analyzing a finite prefix of a **single run** (trace) of the system
- lightweight technique compared to model checking
- naturally applicable to data streaming contexts





# Monitoring: Monitorable Properties

When the monitor reaches a verdict, the latter is definitive.

**Positively monitorable** properties:

- every system satisfying it features a finite trace witnessing the satisfaction
- $\Diamond(ack)$ , at a certain point the system reaches an *ack* state

**Negatively monitorable** properties:

- every system violating it features a finite trace witnessing the violation
- $\Box(online)$ , the system is always *online*

**Not all properties** are monitorable:

- $\Box(req \rightarrow \Diamond(ack))$ , every request submitted to the system ultimately receives an answer



# Linear Temporal Logic [Pnueli 1977]

Linear Temporal Logic (LTL) allows one to express temporal properties over linear structures (single computation paths).

$$\varphi := \top \mid p \mid \neg\varphi \mid \varphi_1 \wedge \varphi_2 \mid \varphi_1 U \varphi_2 \mid X\varphi$$

$$\pi, s_i \models p \in P \Leftrightarrow V(p, s_i) = \text{true}$$

$$\pi, s_i \models \neg\alpha \iff \pi, s_i \not\models \alpha$$

$$\pi, s_i \models \alpha \wedge \beta \iff \pi, s_i \models \alpha \wedge \pi, s_i \models \beta$$

$$\pi, s_i \models X\alpha \iff \pi, s_{i+1} \models \alpha$$


$$\pi, s_i \models F\alpha \iff \exists j \geq i : \pi, s_j \models \alpha$$

$$\pi, s_i \models G\alpha \iff \forall j \geq i : \pi, s_j \models \alpha$$

$$\pi, s_i \models \alpha U \beta \iff \exists j \geq i : \pi, s_j \models \beta \text{ e } \forall k \text{ s.t. } i \leq k < j : \pi, s_k \models \alpha$$

While being very intuitive, LTL cannot handle continuous time series data, but a preliminary discretization step is needed.

SAX (Symbolic Aggregate approXimation) transforms a real-valued time series into a discrete sequence of states



$$G(d \rightarrow F(e))$$





# LTL: Reduction to Finite Model Checking

- *BayesLTL* is a tool for LTL finite model checking and property extraction
- Nevertheless, the solution does not provide any support for monitoring
- Thus, we extended it with such a capability, and we devised a reduction from monitoring to finite model checking
- Since a finite model checking algorithm returns a Boolean answer ( $\top/\perp$ ), while a monitor can also provide an undefined one (?), we first gave a transformation  $\tau$  from an LTL formula  $\varphi$  to a pair of LTL formulas  $\tau(\varphi) = \langle \varphi_1, \varphi_2 \rangle$
- Then we showed that monitoring  $\varphi$  against a given trace amounts to applying the finite model checking algorithm to  $\varphi_1$  and  $\varphi_2$ , and suitably interpreting the outcomes

In what follows, given a pair of formulas  $\tau(\varphi) = \langle \varphi_1, \varphi_2 \rangle$ , we denote by  $\tau(\varphi)_{|_1}$  (resp.,  $\tau(\varphi)_{|_2}$ ) the first (resp., second) formula of the pair, that is,  $\tau(\varphi)_{|_i} = \varphi_i$  ( $i \in \{1, 2\}$ )

## Definition

The mapping  $\tau : LTL \rightarrow LTL \times LTL$  is inductively defined as:

- ①  $\tau(p) = \langle p, p \rangle$ , for all  $p \in \mathcal{AP}$ ;
- ②  $\tau(\neg\psi) = \langle \neg\tau(\psi)_{|_1} \wedge \neg\tau(\psi)_{|_2}, \neg\tau(\psi)_{|_1} \rangle$ ;
- ③  $\tau(\psi \vee \xi) = \langle \tau(\psi)_{|_1} \vee \tau(\xi)_{|_1}, \tau(\psi)_{|_2} \vee \tau(\xi)_{|_2} \rangle$ ;
- ④  $\tau(X\psi) = \langle X\tau(\psi)_{|_1}, X\tau(\psi)_{|_2} \rangle$ ;
- ⑤  $\tau(\psi U \xi) = \langle \tau(\psi)_{|_1} U \tau(\xi)_{|_1}, ((\tau(\psi)_{|_1} \vee \tau(\psi)_{|_2}) U (\tau(\xi)_{|_1} \vee \tau(\xi)_{|_2})) \vee G(\tau(\psi)_{|_1} \vee \tau(\psi)_{|_2}) \rangle$ .

Now, it is possible to reduce the monitoring problem to the finite model checking problem as follows:

## Theorem

Let  $\varphi$  be an LTL formula and  $\pi \in \Sigma^*$  be a finite trace. It holds:

- ①  $\text{monitoring}(\pi, \varphi)$  returns  $\top$  iff  $\pi \models \tau(\varphi)_{|_1}$ ,
- ②  $\text{monitoring}(\pi, \varphi)$  returns  $\perp$  iff  $\pi \not\models \tau(\varphi)_{|_1}$  and  $\pi \not\models \tau(\varphi)_{|_2}$ , and
- ③  $\text{monitoring}(\pi, \varphi)$  returns  $?$  iff  $\pi \not\models \tau(\varphi)_{|_1}$  and  $\pi \models \tau(\varphi)_{|_2}$ .


Given an LTL formula  $\varphi$  and a finite trace  $\pi$ , the monitoring algorithm transforms  $\varphi$  in  $\tau(\varphi) = \langle \varphi_1, \varphi_2 \rangle$ . Then, it runs the *BayesLTL* finite model checking procedure to check  $\varphi_1$  against  $\pi$ : if  $\pi \models \varphi_1$ , then  $\top$  is returned. Otherwise, *BayesLTL* is run again to check  $\varphi_2$  against  $\pi$ : if  $\pi \not\models \varphi_2$ , then  $\perp$  is returned; otherwise, the monitor returns an undefined verdict (?)

Signal Temporal Logic (STL) is an extension of LTL with *real-time* and *real-valued* constraints

$\varphi ::=$	$f(\mathbf{x}) \sim 0$		$f: \mathbb{D} \rightarrow \mathbb{R}$ is a function over the signal $\mathbf{x}: \mathbb{T} \rightarrow \mathbb{D}$ , $\sim \in \{\leq, <, >, \geq, =, \neq\}$
	$\neg \varphi$		Negation
	$\varphi \wedge \varphi$		Conjunction
	$\mathbf{F}_{[a,b]} \varphi$		At some <b>F</b> uture step in the interval $[a, b]$
	$\mathbf{G}_{[a,b]} \varphi$		<b>G</b> lobally in all times in the interval $[a, b]$
	$\varphi \mathbf{U}_{[a,b]} \varphi$		In all steps <b>U</b> ntil in interval $[a, b]$
	$\varphi \mathbf{S}_{[a,b]} \varphi$		In all steps <b>S</b> ince in interval $[a, b]$


$$\mathbf{G}_{[0,3]}((T > 60) \wedge (T < 75))$$

always, between time 0 and time 3,  $60 < T < 75$



$$\mathbf{F}_{[0,60]}(\mathbf{G}(|x| < 0.1))$$

eventually, at some time  $t$  between 0 and 60, from  $t$  on,  $|x| < 0.1$



The satisfaction of a formula  $\varphi$  by a (multivariate) signal  $x = (x_1, \dots, x_n)$  at time  $t$  is given by:

$$\begin{aligned}(\mathbf{x}, t) \models \mu & \Leftrightarrow f(x_1[t], \dots, x_n[t]) > 0 \\(\mathbf{x}, t) \models \varphi \wedge \psi & \Leftrightarrow (x, t) \models \varphi \wedge (x, t) \models \psi \\(\mathbf{x}, t) \models \neg \varphi & \Leftrightarrow \neg((x, t) \models \varphi) \\(\mathbf{x}, t) \models \varphi \mathcal{U}_{[a,b]} \psi & \Leftrightarrow \exists t' \in [t + a, t + b] \text{ such that } (x, t') \models \psi \wedge \\& \quad \forall t'' \in [t, t'], (x, t'') \models \varphi\end{aligned}$$

Note that:

- $\mathbf{F}_{[a,b]} \varphi = \top \mathcal{U}_{[a,b]} \varphi$
- $\mathbf{G}_{[a,b]} \varphi = \neg(\mathbf{F}_{[a,b]} \neg \varphi)$

STL also quantifies the *robustness degree* of satisfaction of a formula by a given trace  $x$  at time  $t$


$$\rho(\top, x, t) = +\infty$$

$$\rho(x_i \geq c, x, t) = x_i(t) - c$$

$$\rho(\neg\phi, x, t) = -\rho(\phi, x, t)$$

$$\rho(\phi_1 \wedge \phi_2, x, t) = \min\{\rho(\phi_1, x, t), \rho(\phi_2, x, t)\}$$

$$\rho(\phi_1 U_I \phi_2, x, t) = \sup_{t_1 \in t+I} \min\{\rho(\phi_2, x, t_1), \inf_{t_2 \in [t, t_1)} \rho(\phi_1, x, t_2)\}$$








# Framework: General Idea

- We consider a pool of monitored properties, that ought to detect or predict system failures in an online fashion
- At each time instant, such properties are checked against the incoming execution trace of the system
- If a failure is detected, the trace is divided into a *good* and a *bad* part, and we look for new properties capable of discerning between such sub-traces
- The new properties are added to the monitoring pool and the monitoring process is resumed
- Intuitively, the framework can be initialized with a very small pool of simple properties
- Then, over time, the pool will be automatically extended with new properties capable of increasing the completeness and preemptiveness of failure detection


# Framework: Monitoring Pool

- The properties in the pool are expressed by means of a suitable temporal logic (in the remainder we focus on STL)
- Actually, in a more general sense, properties can be encoded by a combination of STL formulas, relying on decision tree models (given their interpretability)



# Framework: Failure Detection


- When a failure is predicted by a property in the pool, the incoming trace is divided into a *good* and a *bad* part, according to a windowing approach
- The length of the window is a fixed hyperparameter of the framework



# Framework: Learning of New Properties

## Extraction of new formulas

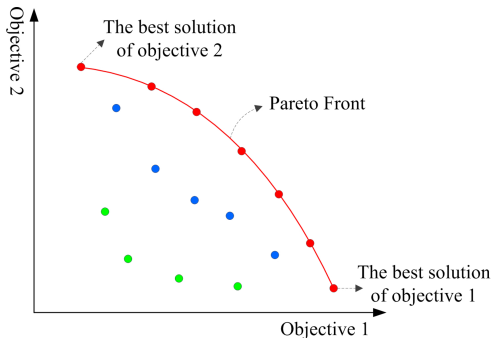
- At this point, new properties are extracted that are capable of discerning between the *good* and the *bad* trace
- A genetic algorithm (GA) is employed that tries to generate highly discriminative and robust STL formulas (2 objs)
- In order to avoid overfitting, starting from the good and bad traces, new traces are generated by applying different kinds of transformation. This is the training set of the GA.



# Framework: Learning of New Properties

## Some thoughts about the extracted formulas

- Since the genetic algorithm follows two (maximization) objectives, a set of optimal solutions (formulas) is produced at the end of its execution (Pareto front)
- Some of them won't be very useful
- Some would be more effective if combined with others





# Framework: Learning of New Properties

## Combination of the extracted formulas

- A dataset is built to support the supervised learning of a decision tree model, where each instance corresponds to a subtrace used during the GA operation
- Each instance is represented by a set of Boolean predictors, one for each extracted formula
- Each predictor is true if and only if the corresponding formula is satisfied by the instance

	$f_1$	$f_2$	$f_3$	...	$f_n$	CASE
Good	T	F	F		T	C
Good'	T	F	T	...	T	C
Good''	T	F	F		F	C
⋮			⋮			⋮
BAD	F	F	T		T	⋮
BAD'	F	T	F	...	F	⋮
BAD''	F	T	T		F	⋮
⋮						⋮



# Framework: Monitoring Pool Update

- The decision tree model is added to the monitoring pool
- Intuitively, such a model will be capable of predicting a forthcoming failure earlier than the property that initially triggered the process that led to its generation
- To each property in the pool, a validity score is attached, that tracks its performance in the detection of failures (F1 score, jointly considering precision and recall measures)
- In this way, the pool is constantly updated: redundant or under-performing properties are removed



# Framework: Source Code

## Algorithm 1 Framework execution

**Input:** initial pool of properties  $\mathcal{P}$ , incoming system trace  $t$

---

```

1: while True do
2:   if  $m \in \mathcal{P}$  predicts a failure in  $t$  at time  $i$ 
     or FAILURE( $i$ ) then
3:     UPDATEPOOLINFORMATION( $\mathcal{P}, t, i$ )
4:      $T \leftarrow$  GENERATETRAINDATA( $t, i$ )
5:      $F \leftarrow$  EXTRACTDISCRFORMULAS( $T$ )
6:      $m' \leftarrow$  BUILDCLASSIFIER( $T, F$ )
7:     CHECKANDADD( $m', \mathcal{P}$ )
8:     SYSTEMFIXANDRESTART()
```

---

## Algorithm 3 BUILDCLASSIFIER

**Input:** training data  $T$ , list of extracted formulas  $F$

---

```

1:  $X \leftarrow$  new empty ( $\text{length}(T) \times \text{length}(F)$ ) matrix
2:  $y \leftarrow$  new empty array of  $\text{length}(T)$  elements
3: for  $t$  from 1 to  $\text{length}(T)$  do
4:   for  $f$  from 1 to  $\text{length}(F)$  do
5:      $X[t][f] \leftarrow$  MONITORING( $T[t], F[f]$ )
6:    $y[t] \leftarrow T[t].\text{label}$ 
7: return TRAINCLASSIFIER( $X, y$ )
```

---

## Algorithm 2 UPDATEPOOLINFORMATION

**Input:** pool of properties  $\mathcal{P}$ , trace  $t$ , failure timestep  $i$

**Global:** forget rate  $\alpha$ , minimum goodness  $g_{\min}$

---

```

1:  $\mathcal{M} \leftarrow$  GETTRIGGEREDCLASSIFIERS( $\mathcal{P}, t, i$ )
2: for  $m \in \mathcal{M}$  do
3:    $\text{good}_m \leftarrow (1 - \alpha) * \text{NEWF1SCORE}(m, t) +$ 
4:      $\alpha * \text{good}_m$ 
5:   if  $\text{good}_m < g_{\min}$  then
6:     REMOVE( $m, \mathcal{P}$ )
7: HANDLEREDUNDANCY( $\mathcal{P}$ )
```

---

## Table 1 Framework hyperparameters

Description	Value	Search range
$\alpha$ forget rate for properties goodness measure update	0.9	{0.7, 0.8, 0.9}
$g_{\min}$ minimum goodness, properties goodness threshold	0.9	{0.7, 0.8, 0.9}
$h_{\max}$ maximum height for property tree representations	3	{2, 3, 4, 5}
$n_{ F }$ number of formulas obtained in the extraction phase	10	{5, 10, 15, 20}
$n_{\text{aug}}$ number of augmentations for each failure trace	100	{50, 100, 150}
$l$ failure window length for generating train data	–	domain specific



## Execution Modes:


- *warmup*: mimic the continual arrival of the available traces from data pertaining to past system failures or generated by means of simulations
- *online*: incoming traces of the currently monitored system are considered

## Execution Strategies:

- *semi-supervised*: domain experts specify an initial set of properties to be monitored
- *unsupervised*: monitor initialized with just a single “the machinery is in operation” property

# Application: Backblaze Hard Drive Dataset

- Information regarding the health status of ST4000DM000 hard drive model in the Backblaze data center
- Data recorded daily from 2015 to 2017
- 21 SMART parameters including both discrete and real values
- Label which indicates a drive failure





# Application: Experiment Setup



- Initial *unsupervised learning warmup* phase performed concatenating a series of training set execution traces
- Two evaluation modes:
  - *offline*, for SOTA comparison purposes
  - *online*, in which the framework continues to learn properties from the execution traces of the test set
- Counter-overfitting measures (trees):
  - maximum height of 3
  - minimum F1 score of 0.9


	S1 LTL	S1 STL	S1 (Huang, 2017, NN)	S2 LTL	S2 STL	S2 (Su, 2019, LSTM)
<b>Precision</b>	0.71	0.73	0.50	0.91	0.97	0.91
<b>Recall</b>	0.43	0.42	0.53	0.85	0.83	0.94
<b>FAR</b>	0.02	0.03	0.01	0.07	0.08	0.05
<b>F<sub>1</sub> score</b>	0.53	0.53	0.52	0.88	0.89	0.93

$$\text{Precision} = \frac{TP}{TP + FP}, \quad \text{Recall} = \frac{TP}{TP + FN},$$

$$\text{FAR} = \frac{FP}{FP + TN}, \quad \text{F1} = \frac{2 * \text{Precision} * \text{Recall}}{\text{Precision} + \text{Recall}}.$$

# Application: RUL and Online Results





The decision tree issues a failure prediction for a hard disk if the latter:

- in the first three days, maintains a temperature exceeding 45.6 °C
- on the third day, also its *uncorrectable sector count* value becomes greater than 0.32, and remains so also for the following day



## Application: Interpretability (2)

Pattern witnessed during the *warmup* phase:

- ① Formula  $f_1 = \mathbf{F}_{[25,45]} SMART_{198} > 2.59$  is extracted at iteration  $i$ 
  - critical sensor regarding *sector read/write errors*
- ② Formula  $f_1$  triggers a failure prediction at iteration  $j > i$
- ③ As a consequence,  $f_2 = \mathbf{F}_{[11,36]} SMART_{189} > 8.28$  is extracted at iteration  $j$ 
  - non-critical sensor regarding *unsafe fly height conditions*

The disk head is operating at an unsafe height, ultimately damaging a disk sector and consequently causing read and write errors (link between a non-critical and a critical sensor).



- Formula-dependent failure windows
- How to estimate RUL for the formulas extracted during the online phase?
- Experimentation with different logic formalism and case studies





- Giuseppe Bombara et al. (2016). “A decision tree approach to data classification using signal temporal logic”. In: *Proceedings of the 19th International Conference on Hybrid Systems: Computation and Control*, pp. 1–10.
- Andrea Brunello, Dario Della Monica, and Angelo Montanari (2019). “Pairing Monitoring with Machine Learning for Smart System Verification and Predictive Maintenance”. In: *Proceedings of the 1st Workshop on Artificial Intelligence and Formal Verification, Logic, Automata, and Synthesis*. Vol. 2509. CEUR Workshop Proceedings, pp. 71–76.
- Andrea Brunello, Dario Della Monica, et al. (2020). “Learning How to Monitor: Pairing Monitoring and Learning for Online System Verification”. In.

- Ian Cassar et al. (2017). “A Survey of Runtime Monitoring Instrumentation Techniques”. In: *Proceedings of the 2nd International Workshop on Pre- Post-Deployment Verification Techniques*, pp. 15–28.
- Alexandre Donzé (2013). “On signal temporal logic”. In: *International Conference on Runtime Verification*. Springer, pp. 382–383.
- Marcus Gerhold, Arnd Hartmanns, and Mariëlle Stoelinga (2019). “Model-Based Testing of Stochastically Timed Systems”. In: *Innovations in Systems and Software Engineering* 15.3-4, pp. 207–233. DOI: 10.1007/s11334-019-00349-z. URL: <https://doi.org/10.1007/s11334-019-00349-z>.
- N. Jansen et al. (2018). “Machine Learning and Model Checking Join Forces”. In: *Dagstuhl Reports* 8.3, pp. 74–93.



- Martin Leucker and Christian Schallhart (2009). “A brief account of Runtime Verification”. In: *Journal of Logical and Algebraic Methods in Programming* 78.5, pp. 293–303. ISSN: 1567-8326. DOI: <http://dx.doi.org/10.1016/j.jlap.2008.08.004>. URL: <http://www.sciencedirect.com/science/article/pii/S1567832608000775>.
- Oded Maler and Dejan Nickovic (2004). “Monitoring temporal properties of continuous signals”. In: *Formal Techniques, Modelling and Analysis of Timed and Fault-Tolerant Systems*. Springer, pp. 152–166.
- Amir Pnueli (1977). “The temporal logic of programs”. In: *Proceedings of the 18th Annual Symposium on Foundations of Computer Science*. IEEE, pp. 46–57.



## Bibliography IV

### *Property Pattern Mappings for LTL (n.d.).*

<https://matthewbdwyer.github.io/psp/>. Kansas State University CIS Department, Laboratory for Specification, Analysis, and Transformation of Software (SAnToS Laboratory) – accessed online on 17 July 2020.